Congratulations! You passed!

TO PASS 80% or higher

Keep Learning

GRADE 100%

Practical aspects of deep learning

LATEST SUBMISSION GRADE

100%

1. If you have 10,000,000 examples, how would you split the train/dev/test set? 1 / 1 point

- 33% train . 33% dev . 33% test
- 60% train . 20% dev . 20% test
- 98% train . 1% dev . 1% test

Correct

2. The dev and test set should: 1 / 1 point

- Come from the same distribution
- Come from different distributions
- Be identical to each other (same (x,y) pairs)
- **Have the same number of examples**

Correct

3. If your Neural Network model seems to have high variance, what of the following would be promising things to try? 1 / 1 point

- [x] Add regularization

Correct

- [ ] Get more test data
- [ ] Increase the number of units in each hidden layer
- [ ] Make the Neural Network deeper
- [x] Get more training data

Correct

4. You are working on an automated check-out kiosk for a supermarket, and are building a classifier for apples, bananas and oranges. Suppose your classifier obtains a training set error of 0.5%, and a dev set error of 7%. Which of the following are promising things to try to improve your classifier? (Check all that apply.) 1 / 1 point

- [x] Increase the regularization parameter lambda

Correct

- [ ] Decrease the regularization parameter lambda
- [x] Get more training data

Correct

- [ ] Use a bigger neural network

5. What is weight decay? 1 / 1 point

- A technique to avoid vanishing gradient by imposing a ceiling on the values of the weights.
- A regularization technique (such as L2 regularization) that results in gradient descent shrinking the weights on every iteration.
- Gradual corruption of the weights in the neural network if it is trained on noisy data.
- The process of gradually decreasing the learning rate during training.

Correct