

Autonomous driving (case study)

LATEST SUBMISSION GRADE
100%

1. To help you practice strategies for machine learning, in this week we'll present another scenario and ask how you would *also* think this "simulator" of working in a machine learning project will give a taste of what leading a machine learning project could be like.

5/11 points

You are employed by a startup building self-driving cars. You are in charge of detecting road signs (stop sign, pedestrian crossing sign, construction ahead sign) and traffic signals (red and green lighting in images). The goal is to recognize where these objects appear in each image. As an example, the above image contains a pedestrian crossing sign and red traffic lights.



$$y^{(1)} = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 1 \\ 0 \end{bmatrix}$$

"stop sign"
"pedestrian crossing sign"
"construction ahead sign"
"red traffic light"
"green traffic light"

Your 100,000 labeled images are taken using the front-facing camera of your car. This is also the distribution of data you can most about during test on. You think you might be able to get a much larger dataset off the internet, that could be helpful for making even if the distribution of internet data is not the same.

You are just getting started on this project. What is the first thing you do? Assume each of the steps below would take about an equal amount of time to few days.

- ☐ Spend a few days checking what is human-level performance for these tasks so that you can get an accurate measure of Bayes error.
- ☐ Spend a few days getting the internet data, so that you understand better what data is available.
- ☒ Spend a few days training a basic model and use what modules it makes.
- ☐ Spend a few days collecting more data using the front-facing camera of your car, to better understand how much data you can collect.

✓ **Correct**

As discussed in lecture, applied ML is a highly iterative process. If you train a basic model and compare out error analysis (see what modules it makes) it will help point you in more promising directions.

2. Your goal is to detect road signs (stop sign, pedestrian crossing sign, construction ahead sign) and traffic signals (red and green lighting in images). The goal is to recognize which of these objects appear in each image. You plan to use a deep neural network with fully connected units in the hidden layers.

5/11 points

For the output layer, a softmax activation would be a good choice for the output layer because this is a multi-class learning problem. True/False?

- ☐ True
- ☒ False

✓ **Correct**

Softmax would be a good choice if one and only one of the possibilities (stop sign, speed limit, pedestrian crossing, green light and red light) was present in each image.

3. You are carrying out error analysis and counting up what error the algorithm makes. What of these datasets do you think you should manually go through and carefully examine, one image at a time?

5/11 points

- ☐ 500 randomly chosen images.
- ☐ 10,000 images on which the algorithm made a mistake.
- ☒ 500 images on which the algorithm made a mistake.
- ☐ 10,000 randomly chosen images.

✓ **Correct**

Scoring on images that the algorithm got wrong, also, 500 is enough to give you a good initial sense of the error statistics. There's probably no need to look at 10,000, which will take a long time.

4. After working on the data for several weeks, your team ends up with the following data:

5/11 points

- 100,000 labeled images taken using the front-facing camera of your car.
- 500,000 labeled images of roads downloaded from the internet.
- Each image's labels precisely indicate the presence of any specific road signs and traffic signals or combinations of them. For example, $y^{(1)} = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 1 \\ 0 \end{bmatrix}$ means the image contains a stop sign and a red traffic light.

Because this is a multi-class learning problem, you need to have all your $y^{(1)}$ entries fully labeled. If one example is equal to $\begin{bmatrix} 0 \\ 1 \\ 0 \\ 1 \\ 0 \end{bmatrix}$, then the learning algorithm will not be able to use that example. True/False?

- ☐ True
- ☒ False

✓ **Correct**

As seen in the lecture on multi-task learning, you can compute the cost such that it is not influenced by the fact that some entries haven't been observed.

5. The distribution of data you can collect contains images from your car's front-facing camera, which comes from a different distribution than the images you were able to find and download off the internet. How should you split the dataset into training and test sets?

5/11 points

- ☐ Mix all the 100,000 images with the 500,000 images you found online. Shuffle everything. Split the 1,000,000 images dataset into 900,000 for the training set, 10,000 for the dev set, and 10,000 for the test set.
- ☐ Mix all the 100,000 images with the 500,000 images you found online. Shuffle everything. Split the 1,000,000 images dataset into 400,000 for the training set, 200,000 for the dev set, and 200,000 for the test set.
- ☐ Choose the training set to be the 500,000 images from the internet along with 20,000 images from your car's front-facing camera. The 50,000 remaining images will be split equally in dev and test sets.
- ☒ Choose the training set to be the 500,000 images from the internet along with 80,000 images from your car's front-facing camera. The 20,000 remaining images will be split equally in dev and test sets.

✓ **Correct**

Yes, as we've discussed, it is important that your dev and test sets have the closest possible distribution to "real" data. It is also important for the training set to contain enough "real" data to avoid having a data-mismatch problem.

6. Assume you've finally chosen the following split between all the data:

5/11 points

Dataset	Contents	Error of the algorithm
Training	100,000 images randomly picked from 500,000 internet images + 50,000 car's front-facing camera images	0.2%
Training-Dev	20,000 images randomly picked from 500,000 internet images + 5,000 car's front-facing camera images	0.1%
Dev	20,000 images from your car's front-facing camera	14.0%
Test	20,000 images from the car's front-facing camera	14.0%

You also know that human performance on the road signs and traffic signals distribution task is around 0.5%. Which of the following are True? Check all that apply.

- ☐ Your algorithm overfits the dev set because the error of the dev and test sets are very close.
- ☒ You have a large data-mismatch problem because your model does a lot better on the training dev set than on the dev set.

✓ **Correct**

You have a large variance problem because your model is not generalizing well to data from the same training distribution but that it has never seen before.

You have a large available bias problem because your training error is quite a bit higher than the human-level error.

✓ **Correct**

You have a large variance problem because your training error is quite higher than the human-level error.

7. Based on data from the previous question, a friend thinks that the training data distribution is much closer than the dev set distribution. What do you think?

5/11 points

- ☐ Your friend is right. (i.e., Bayes error for the training data distribution is probably lower than for the dev set distribution.)
- ☐ Your friend is wrong. (i.e., Bayes error for the training data distribution is probably higher than for the dev set distribution.)
- ☒ There's insufficient information to tell if your friend is right or wrong.

✓ **Correct**

The algorithm does better on the distribution of data it trained on, but you don't know if it's because it trained on that data distribution or if it's really a bias. To get a better sense, measure human-level error separately on both distributions.

8. You decide to focus on the dev set and check by hand what are the errors due to. Here is a table summarizing your discoveries:

5/11 points

Overall dev set error	13.0%
Errors due to incorrectly labeled data	4.1%
Errors due to foggy pictures	0.0%
Errors due to train stops stuck on your car's front-facing camera	0.0%
Errors due to other causes	1.0%

In this table, 4.1%, 0.0%, etc. are a fraction of the total dev set (not just examples your algorithm misclassified. For example, about 0.07% = 0.2% of your errors are due to foggy pictures.

The results from this analysis implies that the team's highest priority should be to bring more foggy pictures into the training set so as to address the 0.0% of errors in that category. True/False?

Additional note: there are subtle concepts to consider with this question, and you may find arguments for why some answers are also correct or incorrect. We recommend that you spend time reading the feedback for this question to understand what issues that you will want to consider when you are building your own machine learning system.

- ☐ True because it is the largest category of errors, we should always prioritize the largest category of error as this will make the best use of the team's time.
- ☐ True because it is greater than the other error categories added together (0.0 + 4.1 + 0.2 + 1.0).
- ☒ False because it depends on how easy it is to add foggy data. If foggy data is very hard and costly to collect, it might not be worth the team's effort.
- ☐ First start with the sources of error that are least costly to fix.

✓ **Correct**

Correct feedback: This is the correct answer. You should consider the tradeoff between the data accessibility and potential improvement of your model learned on this additional data.

9. You can buy a specially designed windshield wiper that help wipe off some of the windings on the front-facing camera. Based on the table from the previous question, which of the following statements do you agree with?

5/11 points

- ☒ 2.2% would be a reasonable estimate of the maximum amount this windshield wiper could improve performance.
- ☐ 2.2% would be a reasonable estimate of the minimum amount this windshield wiper could improve performance.
- ☐ 2.2% would be a reasonable estimate of how much this windshield wiper will improve performance.
- ☐ 2.2% would be a reasonable estimate of how much this windshield wiper could worsen performance in the worst case.

✓ **Correct**

Yes, you will probably not improve performance by more than 2.2% by solving the windshield problem. If your dataset was infinitely big, 2.2% would be a perfect estimate of the improvement you can achieve by purchasing a specially designed windshield wiper that removes the windings.

10. You decide to use data augmentation to address foggy images. You find 1,000 pictures of fog off the internet, and "add" them to clear images to synthesize foggy days. The idea:

5/11 points



Which of the following statements do you agree with?

- ☐ There is little risk of overfitting to the 1,000 pictures of fog as long as you are combining it with a much larger $> 1,000$ of clear/hazy foggy images.
- ☐ Adding synthesized images that look like real foggy pictures (taken from the front-facing camera of your car) to training dataset won't help the model improve because it will introduce available bias.
- ☒ As long as the synthesized fog looks realistic to the human eye, you can be confident that the synthesized data is accurately capturing the distribution of real foggy images (or a subset of it), since human vision is very accurate for the problem you're solving.

✓ **Correct**

Yes, if the synthesized images look realistic, then the model will just see them as if you had added useful data to identify road signs and traffic signals in a foggy weather. I will very likely help.

11. After working further on the problem, you've decided to correct the incorrectly labeled data on the dev set. Which of these statements do you agree with? Check all that apply.

5/11 points

- ☒ You should also correct the incorrectly labeled data in the test set, so that the dev and test sets continue to come from the same distribution.

✓ **Correct**

Yes because you want to make sure that your dev and test sets come from the same distribution for your algorithm to make your team's decision development process is efficient.

- ☐ You should correct incorrectly labeled data in the training set as well so as to avoid your training set now being more different from your dev set.
- ☐ You should not correct the incorrectly labeled data in the test set, so that the dev and test sets continue to come from the same distribution.

- ☒ You do not necessarily need to fix the incorrectly labeled data in the training set, because it's okay for the training set distribution to differ from the dev and test sets. Note that it is important that the dev set and test set have the same distribution.

✓ **Correct**

That deep learning algorithms are quite robust to having slightly different train and dev distributions.

12. To fix your algorithm only recognizes red and green traffic lights. One of your colleagues in the startup is wanting to work on recognizing yellow traffic lights. (After conversations it is an orange light rather than a yellow light, well see the US convention of calling it yellow.) Images containing yellow lights are quite rare, and the dataset have enough data to build a good model. She hopes you can help her out using transfer learning.

5/11 points

What do you tell your colleague?

- ☒ She should try using weights pre-trained on your dataset, and fine-tune further with the yellow light dataset.
- ☐ If she has less than 1,000 images of yellow lights, randomly sample 10,000 images from your dataset and put your other data together. This prevents your dataset from "overfitting" the yellow light dataset.
- ☐ You cannot help her because the distribution of data you have is different from hers, and is also lacking the yellow light.
- ☐ Recommend that she try multi-task learning instead of transfer learning using all the data.

✓ **Correct**

Yes, you have trained your model on a huge dataset, and she has a small dataset. Although your labels are different, the parameters of your model have been trained to recognize many characteristics of road and traffic images which will be useful for her problem. This is a perfect case for transfer learning, she can start with a model with the same architecture as yours, change what is after the last hidden layer and initialize it with your learned parameters.

13. Another colleague wants to use microphones placed outside the car to better hear if there are other vehicles around you. For example, if there is a police vehicle behind you, you would be able to hear their siren. However, they don't have much to train the audio system. How can you help?

5/11 points

- ☐ Transfer learning from your vision dataset could help your colleague get going faster. Multi-task learning seems significantly more promising.
- ☐ Multi-task learning from your vision dataset could help your colleague get going faster. Transfer learning seems significantly more promising.
- ☐ Either transfer learning or multi-task learning could help your colleague get going faster.
- ☒ Random transfer learning or multi-task learning seems promising.

✓ **Correct**

Yes, the problem he is trying to solve is quite different from yours, the different dataset structures make it probably impossible to use transfer learning or multi-task learning.

14. To recognize red and green lights, you have been using this approach:

5/11 points

- (A) Input an image to a neural network and have it directly learn a mapping to make a prediction as to whether there is a red light and/or green light on.
- (B) Learn to propose a different, two-step approach:

- (A) In this two-step approach, you first do detect the traffic light in the image (if any), then (B) determine the color of the illuminated lamp in the traffic light.

Between these two, approach B is more of an end-to-end approach because it has distinct steps for the input and the output and, True/False?

- ☐ True
- ☒ False

✓ **Correct**

Yes, (A) is an end-to-end approach as it maps directly the input to the output (A).

15. Approach A in the question above tends to be more promising than approach B if you have a _____ in the labels.

5/11 points

- ☒ Large training set
- ☐ Multi-task learning problem.
- ☐ Large bias problem.
- ☐ Problem with a high Bayes error.

✓ **Correct**

Yes, in many fields, it has been observed that end-to-end learning works better in practice, but requires a large amount of data.